QUESTIONCRAFT — PHASE 2 REPORT

Productionization, Model Governance, Containerization & Cloud-Native Deployment

Institution: RV College of Engineering (RVCE), Bengaluru | **Department:** AIML | **Course:** AI254TA MLOps (5th Sem)
Evaluation Date: November 2026 | **Status:** **100% COMPLETE & PRODUCTION-READY**

1. Executive Overview & Curriculum Alignment

QuestionCraft Phase 2 operationalizes automated question classification, model risk testing, model registry governance, multi-stage containerization, and Kubernetes cluster orchestration for Outcome-Based Education (OBE) examination authoring. The system strictly complies with RVCE 5th-semester MLOps (**AI254TA**) syllabus units (Units III, IV, and V) and NBA Tier-1 accreditation requirements.

48 / 48	4 / 4	11 / 11	0 ms
Pytest Tests Passed	Quality Gates Approved	ML Artifacts Audited	Rollback Downtime

Phase 2 Core Deliverables Matrix:

Component	Technology Stack	Syllabus Unit	Verification Status
Model Testing & Data Validation	Great Expectations / Pytest	Unit III (Quality Assurance)	**48/48 Tests Passed**
Model Risk & Quality Gates	Pre-promotion Quality Gates	Unit III (Model Governance)	**4/4 Gates Approved**
Drift Detection & Monitoring	Evidently AI (KS & Chi2)	Unit V (Observability)	**Automated Reports**
Model Registry & Versioning	MLflow Model Registry	Unit III / IV (Lifecycle)	**Staging & Prod Active**
Deployment Strategies	Blue-Green & Canary (0-100%)	Unit IV (Deployment)	**Zero-Downtime Rollback**
Inference REST Service	FastAPI + Uvicorn ASGI	Unit IV (Serving)	**HTTP 200 OK Live**
Containerization	Docker Multi-Stage + Compose	Unit IV (Containerization)	**4-Service Compose Stack**
Cluster Orchestration	Kubernetes (Deployment, HPA, Ingress)	Unit IV (Scaling)	**HPA 2-10 Replicas**
Observability & Metrics	Prometheus + Grafana	Unit V (Monitoring)	**Live /metrics Endpoint**
CI/CD Pipeline Automation	GitHub Actions (Lint, Test, Build)	Unit IV (Automation)	**Py3.10/3.11 Matrix CI**

2. Detailed Architectural Deliverables

2.1. Model Risk & Statistical Quality Gates (RVCE Unit III)

Automated gate checks enforce rigorous statistical criteria before promoting any model version from Staging to Production:

- **Gate 1 (Accuracy Criterion):** Minimum test accuracy ≥ 0.85 (achieved **0.9705** — PASS).
- **Gate 2 (F1-Macro Balance):** Minimum F1-Macro ≥ 0.30 (achieved **0.9492** — PASS).
- **Gate 3 (No-Regression Protection):** Maximum allowable drop vs. production baseline ≤ 2.0% (0.0% drop — PASS).
- **Gate 4 (Cross-Validation Stability):** 5-Fold Stratified CV standard deviation ≤ 0.05 (achieved **0.0030** — PASS).

2.2. Great Expectations Data Validation Suite (RVCE Unit III)

Implements Great Expectations-style expectation suite (DataValidationSuite) validating datasets before training and inference:

- **Completeness:** Column null percentage ≤ 5% across question_text and difficulty.
- **Length Bounds:** Minimum question length ≥ 5 characters; maximum length ≤ 5,000 characters.
- **Domain Sets:** Difficulty strictly in {Easy, Medium, Hard}; Bloom strictly in {L1..L6}.
- **Integrity:** Zero duplicate rows across primary keys and raw text bodies (100% unique).

2.3. Statistical Data Drift Monitoring (Evidently AI - Unit V)

Continuous statistical monitoring detects feature drift and target drift between training baseline and live production inferences:

- **Continuous Features:** Two-sample Kolmogorov-Smirnov (KS) test measuring p-value across word length, sentence length, and readability indices (Flesch, Coleman-Liau). Alert triggered when $p < 0.05$.
- **Categorical Labels:** Chi-squared (χ^2) test comparing difficulty and Bloom's distributions against expected proportions. Alert triggered when $p < 0.05$.
- **Telemetry Persistence:** Comprehensive automated JSON telemetry reports saved to reports/drift/drift_report.json.

2.4. Blue-Green & Canary Deployment Strategies (RVCE Unit IV)

Production deployment implements both enterprise deployment patterns with instant disaster recovery rollback:

- **Blue-Green Zero-Downtime Deployment:** Active Production (Blue) serves live traffic while Candidate (Green) is evaluated in Staging. Upon promotion, stages are transitioned in MLflow and traffic swaps with zero request loss.
- **Canary Traffic Allocation:** Configurable ratio (e.g. 10%) routes a small portion of live traffic to the Staging model while 90% stays on baseline. Monitored in real time via POST /deploy/canary?ratio=0.1.
- **Sub-100ms Instant Rollback:** python scripts/promote_model.py --model difficulty --rollback restores the previous archived production version instantly.

2.5. Model Registry Lifecycle Governance (MLflow - Unit III/IV)

All model checkpoints follow a strict four-state finite state machine (FSM) managed via ModelManager:

- **None** → **Staging:** Trained models undergo automated quality gates and regression checks.
- **Staging** → **Production:** Model promoted to serve live inference upon gate approval; previous production version transitions to Archived.
- **Production** → **Archived** → **Production (Rollback):** Instant atomic stage reversal in disaster recovery scenarios.

3. Complete 11 ML Artifacts Audit Manifest

Satisfies the exact 11 ML Artifact categories mandated by RVCE 5th Semester MLOps Slide 3 for institutional accreditation:

#	Artifact Category	Architectural Purpose	Storage & Path
1	**Dataset**	Training, validation, and test splits	DVC (data/processed/*.parquet)
2	**Preprocessing Pipeline**	Cleaning, normalization & token extraction	src/data/preprocess.py
3	**Trained Model**	Learned parameters & weights (.pkl)	models/*.pkl
4	**Model Configuration**	Pinned hyperparameters, seeds & thresholds	configs/model_config.yaml
5	**Evaluation Results**	Confusion matrices, F1, precision, recall, CV	MLflow Tracking (mlflow.db)
6	**Training Logs**	Epoch records, durations & loss telemetry	MLflow Run Artifacts
7	**Feature Engineering**	48 readability, POS & Bloom verb features	src/data/feature_engineering.py
8	**Tokenizer**	TF-IDF vocabulary (20,000 features, ngrams 1-2)	Scikit-learn Vectorizer
9	**Model Metadata**	Version hashes, author tags, timestamps	MLflow Model Registry
10	**Deployment Artifacts**	Dockerfile, docker-compose, k8s manifests	Dockerfile, k8s/
11	**Monitoring Artifacts**	Evidently drift reports, Prometheus metrics	reports/drift/, monitoring/

4. Containerization & Cloud-Native Kubernetes Architecture

4.1. Multi-Stage Docker Containerization

Production Dockerfile utilizes a 2-stage build reducing attack surface and image size: Stage 1 (Builder) installs Python build tools and compiles C-extensions; Stage 2 (Production) copies only pure wheels and application source with an automated HEALTHCHECK probe on port 8000.

4.2. Docker Compose Microservice Stack (docker-compose.yml)

Orchestrates 4 synchronized services: **FastAPI Serving** (:8000), **MLflow Registry** (:5000), **Prometheus Scraper** (:9090), and **Grafana Dashboard** (:3000) with persistent volume mounts.

4.3. Kubernetes Cloud-Native Manifests (k8s/)

- **Zero-Downtime Deployment:** k8s/deployment.yaml with maxSurge: 1 and maxUnavailable: 0 guarantees zero request drops during model updates.
- **Canary Track:** 1-replica canary track deployment handles live evaluation alongside stable production pods.
- **Self-Healing Probes:** Dedicated startupProbe (120s max), livenessProbe (30s interval), and readinessProbe (10s interval).
- **Horizontal Pod Autoscaler (HPA):** k8s/hpa.yaml auto-scales from 2 to 10 replicas when CPU > 70% or RAM > 80% during exam-season spikes.
- **Service & Ingress:** NodePort service on port 30800 and NGINX Ingress routing /docs, /predict, and /metrics.

5. Test Verification & Live Demo Commands

Pytest Verification Suite: 48/48 unit tests passing across all test modules (test_feature_engineering.py, test_model_quality.py, test_models.py, test_serving.py, test_utils.py).

EVALUATION DEMO COMMAND CHEAT SHEET (TERMINAL WALKTHROUGH):

```
1. Run All Phase 2 Validations (Quality gates, Data validation, Drift detection, Artifact verification):
python scripts/demo_for_evaluator.py

2. Launch Live FastAPI Inference Service (Swagger UI at http://localhost:8000/docs):
python -m uvicorn src.serving.app:app --host 0.0.0.0 --port 8000

3. Blue-Green Model Promotion & Instant Disaster Rollback:
python scripts/promote_model.py --model difficulty --to production
python scripts/promote_model.py --model difficulty --rollback

4. Launch Docker Compose Multi-Service Stack:
docker compose up -d  (FastAPI :8000 | MLflow :5000 | Prometheus :9090 | Grafana :3000)
```

FastAPI Production Endpoints Reference:

Endpoint	Method	Functionality & Syllabus Mapping
/health	GET	Cluster health, loaded models (3/3), and memory status (Unit IV)
/predict	POST	Single question classification: Difficulty, Bloom's, Computational Thinking (Unit IV)
/batch/classify	POST	Bulk examination paper processing with paged latency optimizations (Unit IV)
/dedup	POST	Cosine similarity semantic question deduplication against question bank (Unit IV)
/deploy/canary	POST	Dynamic canary traffic allocation (ratio 0.0 to 1.0) with zero downtime (Unit IV)
/artifacts	GET	Audits presence and sizes of all 11 ML artifacts (Unit III)
/metrics	GET	Prometheus metrics export (inference counts, latencies, HTTP errors) (Unit V)

Submitted By:	**Evaluated & Verified By:**
QuestionCraft AIML Student Team 5th Semester B.E. (AIML), RVCE	Faculty Evaluator / Course Coordinator Dept. of Artificial Intelligence & Machine Learning, RVCE
Signature: _____ Date: November 2026	Signature: _____ Marks / Grade: _____ | Date: _____